Sutskever, and Dario Amodei. 2020. Language Models are Few-Shot Learners. In *Advances in Neural Information Processing Systems*.

[5] Randal E. Bryant, Randy H. Katz, and Edward D. Lazowska. 2008. Big-Data Computing: Creating Revolutionary Breakthroughs in Commerce, Science, and Society. (2008). https://cra.org/ccc/wp-content/uploads/sites/2/2015/05/Big_Data.pdf Accessed: 2025-03-01.

[6] Sébastien Bubeck, Varun Chandrasekaran, Ronen Eldan, Johannes Gehrke, Eric Horvitz, Ece Kamar, Peter Lee, Yin Tat Lee, Yuanzhi Li, Scott M. Lundberg, Harsha Nori, Hamid Palangi, Marco Túlio Ribeiro, and Yi Zhang. 2023. Sparks of Artificial General Intelligence: Early experiments with GPT-4. *arXiv preprint* arXiv:2303.12712 (2023).

[7] D. Carrell and R. Smith. 2024. Legal AI: Applications and Risks of Large Language Models in Law. *Harvard Law Review* 137 (2024), 112–130. https://www.reuters.com/technology/artificial-intelligence/ai-hallucinations-court-papers-spell-trouble-lawyers-2025-02-18/.

[8] M. et al. Chen. 2021. Evaluating Large Language Models for Code Generation. *arXiv preprint arXiv:2107.03374* (2021). https://arxiv.org/abs/2107.03374.

[9] Yann Collet. 2013. Finite State Entropy. https://github.com/Cyan4973/FiniteStateEntropy. Accessed: 2025-03-01.

[10] Y. Collet. 2016. Zstd github repository from facebook. https://github.com/facebook/zstd

[11] Transaction Processing Performance Council. 2011. TPC-H. http://www.tpc.org/tpch/default.asp

[12] D. Cox. 2016. Syntactically informed text compression with recurrent neural networks. *arXiv preprint* arXiv:1608.02893 (2016).

[13] DeepSeek-AI, Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shirong Ma, Peiyi Wang, Xiao Bi, Xiaokang Zhang, Xingkai Yu, Yu Wu, Z. F. Wu, Zhibin Gou, Zhihong Shao, Zhuoshu Li, Ziyi Gao, Aixin Liu, Bing Xue, Bingxuan Wang, Bochao Wu, Bei Feng, Chengda Lu, Chenggang Zhao, Chengqi Deng, Chenyu Zhang, Chong Ruan, Damai Dai, Deli Chen, Dongjie Ji, Erhang Li, Fangyun Lin, Fucong Dai, Fuli Luo, Guangbo Hao, Guanting Chen, Guowei Li, H. Zhang, Han Bao, Hanwei Xu, Haocheng Wang, Honghui Ding, Huajian Xin, Huazuo Gao, Hui Qu, Hui Li, Jianzhong Guo, Jiashi Li, Jiawei Wang, Jingchang Chen, Jingyang Yuan, Junjie Qiu, Junlong Li, J. L. Cai, Jiaqi Ni, Jian Liang, Jin Chen, Kai Dong, Kai Hu, Kaige Gao, Kang Guan, Kexin Huang, Kuai Yu, Lean Wang, Lecong Zhang, Liang Zhao, Litong Wang, Liyue Zhang, Lei Xu, Leyi Xia, Mingchuan Zhang, Minghua Zhang, Minghui Tang, Meng Li, Miaojun Wang, Mingming Li, Ning Tian, Panpan Huang, Peng Zhang, Qiancheng Wang, Qinyu Chen, Qiushi Du, Ruiqi Ge, Ruisong Zhang, Ruizhe Pan, Runji Wang, R. J. Chen, R. L. Jin, Ruyi Chen, Shanghao Lu, Shangyan Zhou, Shanhuang Chen, Shengfeng Ye, Shiyu Wang, Shuiping Yu, Shunfeng Zhou, Shuting Pan, and S. S. Li. 2025. DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning. (2025). arXiv:2501.12948

[14] G. Delétang, A. Ruoss, P. Duquenne, E. Catt, T. Genewein, C. Mattern, J. Grau-Moya, L. K. Wenliang, M. Aitchison, L. Orseau, M. Hutter, and J. Veness. 2024. Language modeling is compression. In *International Conference on Learning Representations (ICLR)*.

[15] Peter Deutsch. [n. d.]. GZIP file format specification version 4.3. *RFC* 1952 ([n. d.]).

[16] Dmitry Lepikhin et al. 2021. GShard: Scaling Giant Models with Conditional Computation and Automatic Sharding. In *Proceedings of NeurIPS*. https://arxiv.org/abs/2006.16668

[17] Angela Fan, Mike Lewis, and Yann Dauphin. 2018. Hierarchical Neural Story Generation. In *Proceedings of ACL*. https://arxiv.org/abs/1805.04833

[18] Markus Freitag and Yaser Al-Onaizan. 2017. Beam search strategies for neural machine translation. *arXiv preprint arXiv:1702.01806* (2017).

[19] M. Goyal, K. Tatwawadi, S. Chandak, and I. Ochoa. 2019. Deepzip: Lossless data compression using recurrent neural networks. In *Data Compression Conference (DCC)*.

[20] Andrew Gray. 2023. ChatGPT "contamination": estimating the prevalence of LLMs in the scholarly literature. https://en.wikipedia.org/wiki/Generative_artificial_intelligence.

[21] Daya Guo, Qihao Zhu, Dejian Yang, Zhenda Xie, Kai Dong, Wentao Zhang, Guanting Chen, Xiao Bi, Y. Wu, Y.K. Li, Fuli Luo, Yingfei Xiong, and Wenfeng Liang. 2024. DeepSeek-Coder: When the Large Language Model Meets Programming – The Rise of Code Intelligence. https://arxiv.org/abs/2401.14196

[22] Ari Holtzman, Jan Buys, Li Du, Maxwell Forbes, and Yejin Choi. 2020. The Curious Case of Neural Text Degeneration. In *Proceedings of ICLR*. https://arxiv.org/abs/1904.09751

[23] Y. Huang, J. Zhang, Z. Shan, and J. He. 2024. Compression represents intelligence linearly. *arXiv preprint* arXiv:2404.09937 (2024).

[24] D. A. Huffman. 1952. *A Method for the Construction of Minimum-Redundancy Codes*. Vol. 40. 1098–1101 pages. doi:10.1109/JRPROC.1952.273884 Accessed: 2024-11-07.

[25] Binyuan Hui, Jian Yang, Zeyu Cui, Jiaxi Yang, Dayiheng Liu, Lei Zhang, Tianyu Liu, Jiajun Zhang, Bowen Yu, Kai Dang, et al. 2024. Qwen2. 5-Coder Technical Report. *arXiv preprint arXiv:2409.12186* (2024).

[26] IBM Corporation. 2013. Bringing Big Data and Business Analytics to the Enterprise. (2013). https://public.dhe.ibm.com/software/os/systemz/pdf/Bringing_Big_Data_and_Business_Analytics_to_the_Enterprise.pdf Accessed: 2025-03-01.

[27] Robert A. Jacobs, Michael I. Jordan, Steven J. Nowlan, and Geoffrey E. Hinton. 1991. Adaptive Mixtures of Local Experts. *Neural Computation* 3, 1 (03 1991), 79–87.

[28] Prabha Kannan. 2023. How Much Research Is Being Written by Large Language Models? https://en.wikipedia.org/wiki/Generative_artificial_intelligence.

[29] Diederik P. Kingma and Jimmy Ba. 2015. Adam: A Method for Stochastic Optimization. In *International Conference on Learning Representations*.

[30] Sunjun Kweon, Junu Kim, Jiyoun Kim, Sujeong Im, Eunbyeol Cho, Seongsu Bae, Jungwoo Oh, Gyubok Lee, Jong Hak Moon, Seng Chan You, Seungjin Baek, Chang Hoon Han, Yoon Bin Jung, Yohan Jo, and Edward Choi. 2023. Publicly Shareable Clinical Large Language Model Built on Synthetic Clinical Notes. arXiv:2309.00237 [cs.CL]

[31] Guohao Li, Hasan Abed Al Kader Hammoud, Hani Itani, Dmitrii Khizbullin, and Bernard Ghanem. 2023. CAMEL: Communicative Agents for "Mind" Exploration of Large Scale Language Model Society. arXiv:2303.17760 [cs.AI]

[32] Zhenghao Lin, Zhibin Gou, Yeyun Gong, Xiao Liu, Yelong Shen, Ruochen Xu, Chen Lin, Yujiu Yang, Jian Jiao, Nan Duan, and Weizhu Chen. 2024. Rho-1: Not All Tokens Are What You Need. arXiv:2404.07965 [cs.CL]

[33] Jiang Liu, Jialian Wu, Prakamya Mishra, Zicheng Liu, Sudhanshu Ranjan, Pratik Prabhanjan Brahma, Yusheng Su, Gowtham Ramesh, Peng Sun, Zhe Li, Dong Li, Lu Tian, and Emad Barsoum. 2024. AMD-OLMo: A series of 1B language models trained from scratch by AMD on AMD Instinct™ MI250 GPUs. https://huggingface.co/amd/AMD-OLMo

[34] Yu Mao, Yufei Cui, Tei-Wei Kuo, and Chun Jason Xue. 2022. Trace: A fast transformer-based general-purpose lossless compressor. In *Proceedings of the ACM Web Conference 2022*. 1829–1838.

[35] Yu Mao, Jingzong Li, Yufei Cui, and Jason Chun Xue. 2023. Faster and stronger lossless compression with optimized autoregressive framework. In *2023 60th ACM/IEEE Design Automation Conference (DAC)*. IEEE, 1–6.

[36] Sachin Mehta, Mohammad Hossein Sekhavat, Qingqing Cao, Maxwell Horton, Yanzi Jin, Chenfan Sun, Iman Mirzadeh, Mahyar Najibi, Dmitry Belenko, Peter Zatloukal, and Mohammad Rastegari. 2024. OpenELM: An Efficient Language Model Family with Open Training and Inference Framework. *arXiv.org* (April 2024). https://arxiv.org/abs/2404.14619v1

[37] Meta. 2024. Introducing Llama 3.1: Our most capable models to date. https://ai.meta.com/blog/meta-llama-3-1/.

[38] Meta. 2024. Llama 3.2: Revolutionizing edge AI and vision with open, customizable models. https://ai.meta.com/blog/llama-3-2-connect-2024-vision-edge-mobile-devices/.

[39] Arindam Mitra, Hamed Khanpour, Corby Rosset, and Ahmed Awadallah. 2024. Orca-Math: Unlocking the potential of SLMs in Grade School Math. arXiv:2402.14830 [cs.CL]

[40] F. Mittu, Y. Bu, A. Gupta, A. Devireddy, A. E. Ozdarendeli, A. Singh, and G. Anumanchipalli. 2024. Finezip: Pushing the limits of large language models for practical lossless text compression. *arXiv preprint* arXiv:2409.17141 (2024).

[41] Andreas Bentzen Winje Nicolai Thorer Sivesind. 2023. Human-vs-Machine.

[42] OpenAI. 2022. ChatGPT. https://chat.openai.com/chat/.

[43] OpenAI. 2023. ChatGPT. https://en.wikipedia.org/wiki/ChatGPT.

[44] OpenAI. 2023. GPT-4 Technical Report. arXiv:2303.08774 [cs.CL]

[45] Claude Elwood Shannon. 1948. A mathematical theory of communication. *The Bell system technical journal* 27, 3 (1948), 379–423.

[46] Qwen Team. 2024. Qwen2.5: A Party of Foundation Models. https://qwenlm.github.io/blog/qwen2.5/

[47] Brian Thompson, Mehak Dhaliwal, Peter Frisch, Tobias Domhan, and Marcello Federico. 2024. A Shocking Amount of the Web is Machine Translated: Insights from Multi-Way Parallelism. https://en.wikipedia.org/wiki/Generative_artificial_intelligence.

[48] Hugo Touvron, Albert Jiang, et al. 2023. LLaMA: Open and Efficient Foundation Language Models. *arXiv preprint* (2023). https://arxiv.org/abs/2302.13971

[49] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. 2017. Attention is All you Need. In *Advances in Neural Information Processing Systems*. 5998–6008.

[50] Y. Wang and J. Lee. 2023. Applications of Large Language Models in Medical Diagnostics and Education. *Journal of Medical AI Research* 45 (2023), 23–35. https://pmc.ncbi.nlm.nih.gov/articles/PMC11396764/.

[51] Witten et al. 1987. Arithmetic coding for data compression. *Commun. ACM* 30, 6 (1987), 520–540.

[52] Jinzhao Xiao, Yuxiang Huang, Changyu Hu, Shaoxu Song, Xiangdong Huang, and Jianmin Wang. 2022. Time series data encoding for efficient storage: a comparative analysis in Apache IoTDB. *Proc. VLDB Endow.* 15, 10 (June 2022), 2148–2160. doi:10.14778/3547305.3547319

[53] Li Xu, Yanqing He, Guanghui Li, Wei Li, and Wen Shi. 2021. Big Data-Driven Smart Manufacturing: Challenges, Advances and Future Trends. *IEEE Transactions on Industrial Informatics* 17, 3 (2021), 2274–2284.

[54] Vishaal Y. 2025. Synthetic Code Generations Dataset. https://huggingface.co/datasets/VishaalY/synthetic-code-generations.
[55] An Yang, Beichen Zhang, Binyuan Hui, Bofei Gao, Bowen Yu, Chengpeng Li, Dayiheng Liu, Jianhong Tu, Jingren Zhou, Junyang Lin, Keming Lu, Mingfeng Xue, Runji Lin, Tianyu Liu, Xingzhang Ren, and Zhenru Zhang. 2024. Qwen2.5-Math Technical Report: Toward Mathematical Expert Model via Self-Improvement. *arXiv preprint arXiv:2409.12122* (2024).
[56] Hailin Zhang, Penghao Zhao, Xupeng Miao, Yingxia Shao, Zirui Liu, Tong Yang, and Bin Cui. 2023. Experimental Analysis of Large-Scale Learnable Vector Storage Compression. *Proc. VLDB Endow.* 17, 4 (Dec. 2023), 808–822. doi:10.14778/3636218.3636234
[57] Lixi Zhou, Jiaqing Chen, Amitabh Das, Hong Min, Lei Yu, Ming Zhao, and Jia Zou. 2022. Serving deep learning models with deduplication from relational databases. 15, 10 (June 2022), 2230–2243. doi:10.14778/3547305.3547325
[58] Ziv et al. 1977. A universal algorithm for sequential data compression. *IEEE Transactions on information theory* 23, 3 (1977), 337–343.